

Types of Parallel Computers

Flynn's Classical

Taxonomy

S I S D	S I M D
Single Instruction, Single Data	Single Instruction, Multiple Data
M I S D	M I M D
Multiple Instruction, Single Data	Multiple Instruction, Multiple Data

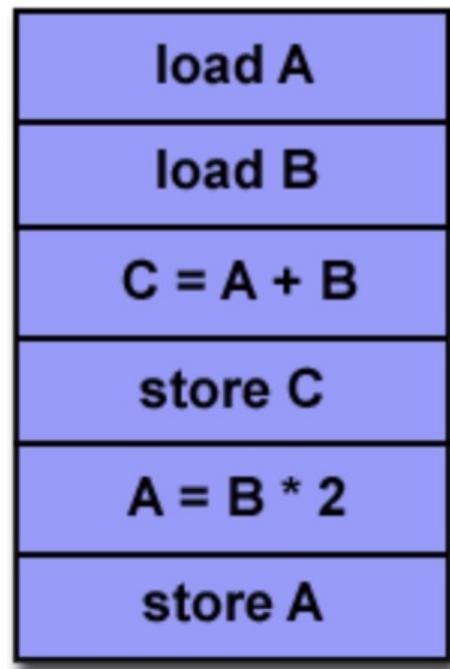
Instruction Set
Data Stream

Parallel

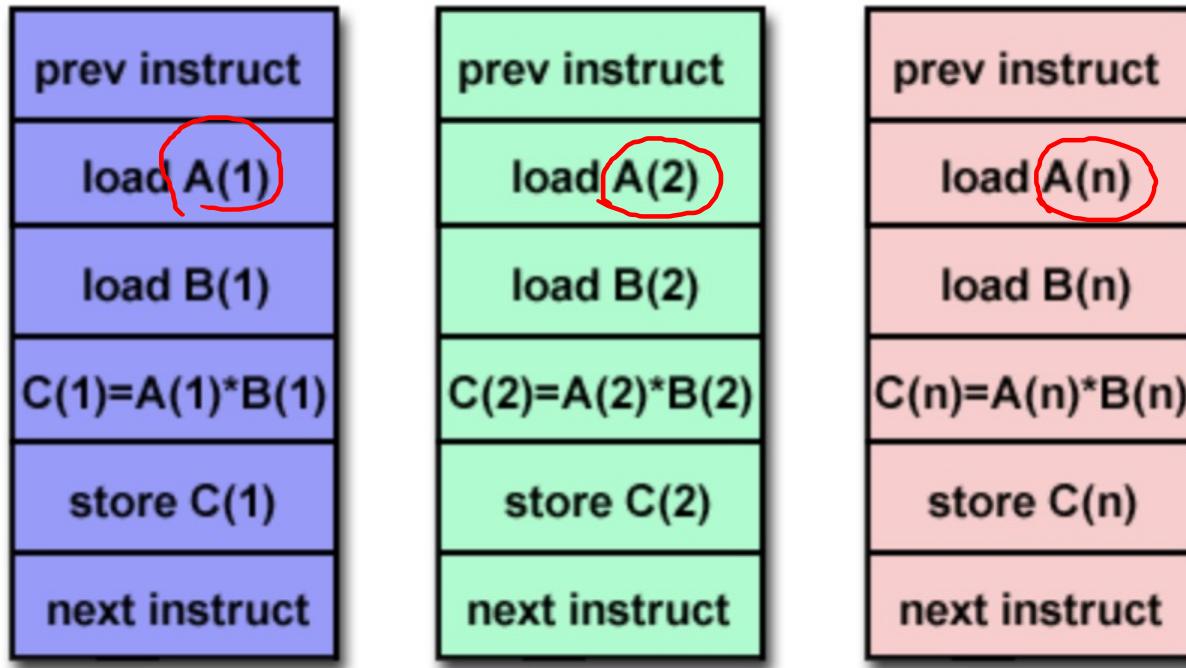
Serial Computer

S

SISD



SIMD



time

MISD



Research



prev instruct
load A(1)
$C(1)=A(1)*1$
store C(1)
next instruct

P1

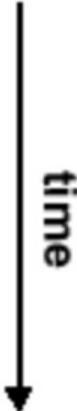
prev instruct
load A(1)
$C(2)=A(1)*2$
store C(2)
next instruct

P2

prev instruct
load A(1)
$C(n)=A(1)*n$
store C(n)
next instruct

Pn

time



MIMD
==

prev instruct
load A(1)
load B(1)
$C(1)=A(1)*B(1)$
store C(1)
next instruct

P1

prev instruct
call funcD
$x=y^z$
$sum=x^2$
call sub1(i,j)
next instruct

P2

prev instruct
do 10 i=1,N
$\alpha=w^{*3}$
$\zeta=C(i)$
10 continue
next instruct

Pn

time

m/y Architectures.

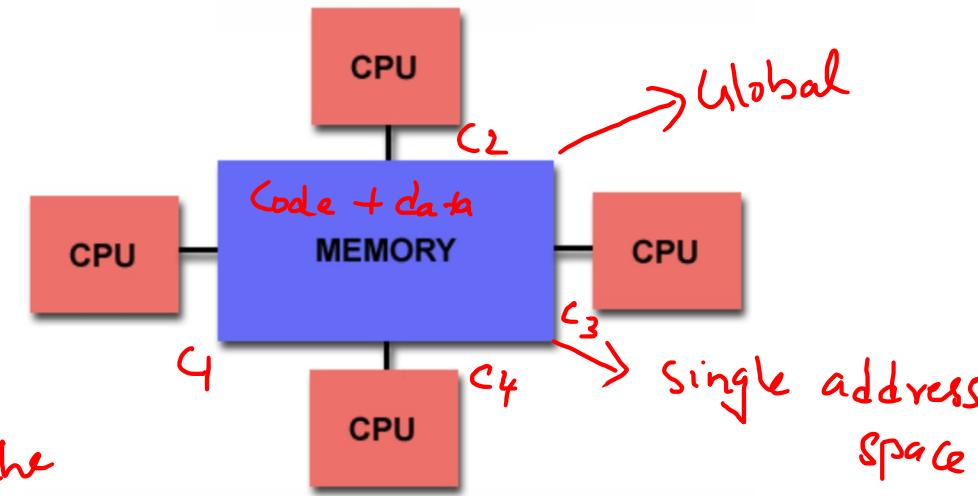
1. Shared

Executable code

Two forms

(1) Uniform m/y access (UMA)

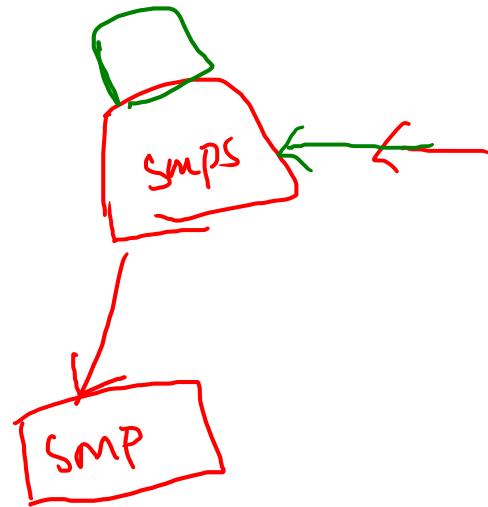
(2) non-uniform m/y access (NUMA)



Cache
preserving
locality

CC - UMA

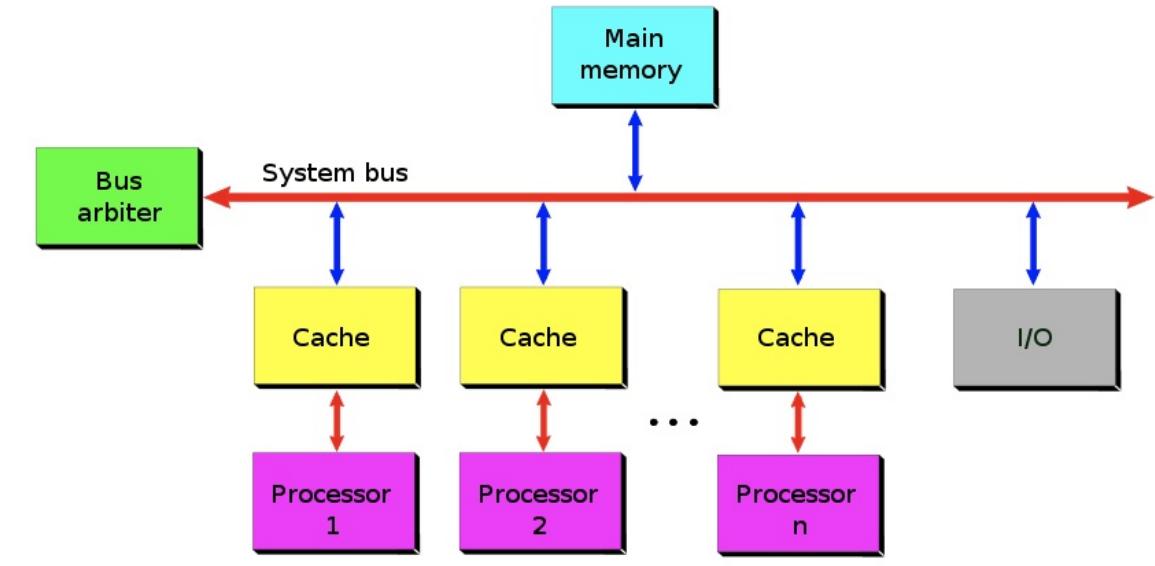
Cache coherent



a processor can
access its own local
memory faster than
non-local memory

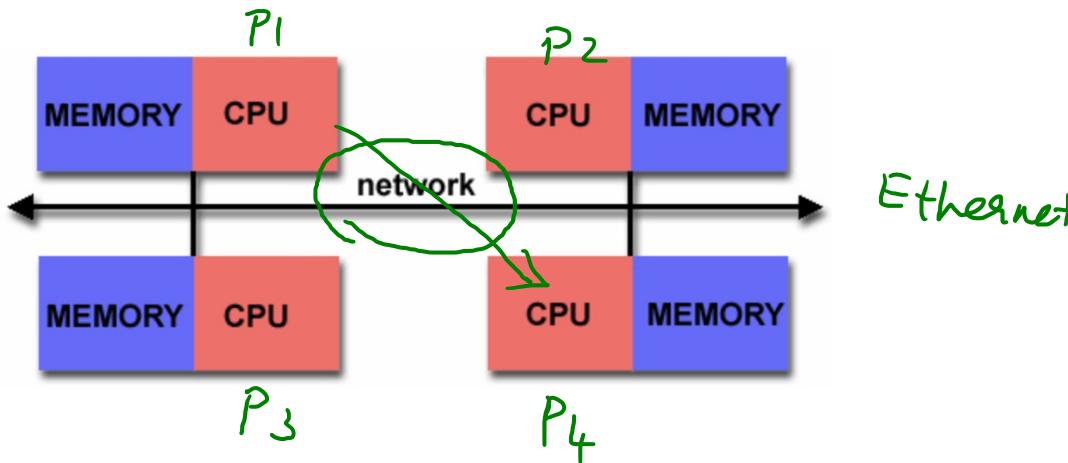
SMP - symmetric multiprocessor system

=



Distributed Memory

my address
↳ no concept
of global address
space.



Concept of cache coherency ? {not applicable}

message passing → Programming

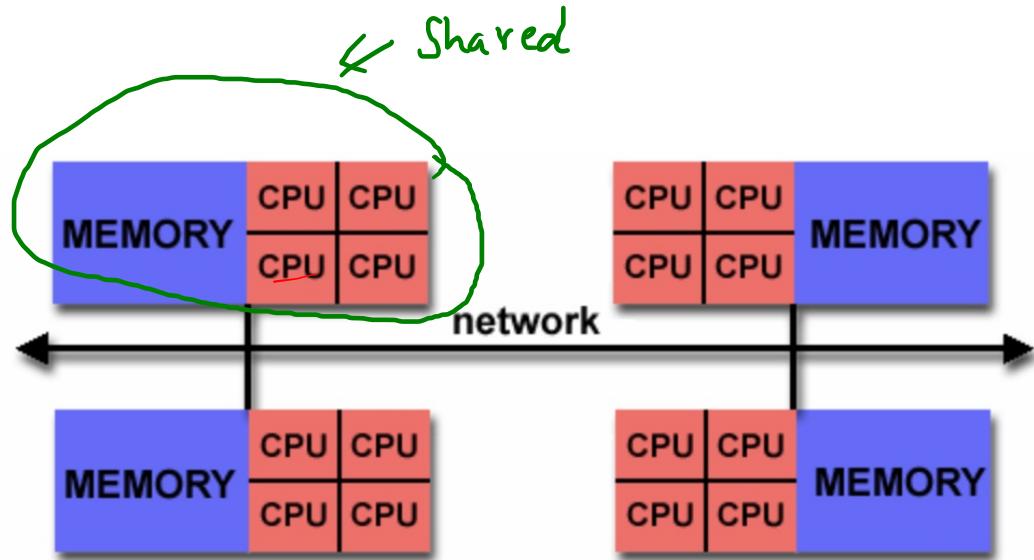
(MPI) =

Hybrid models

=====

Shared +
distributed

Types of parallel
computers



1. Shared
2. Distributed
3. Hybrid

Parallel Programming Models & Paradigms

Traditional — sequential (cost of components)
↓ high availability

parallel paradigms



{cluster} → cost-effective

n/w/s



Super computers

Scalability

Scalable computing clusters ?



{ PC or WS or SMPs. }

homogeneous
heterogeneous

Intelligent mechanisms

(huge data set)

n/w wide resource sharing

↳ Scalability performances

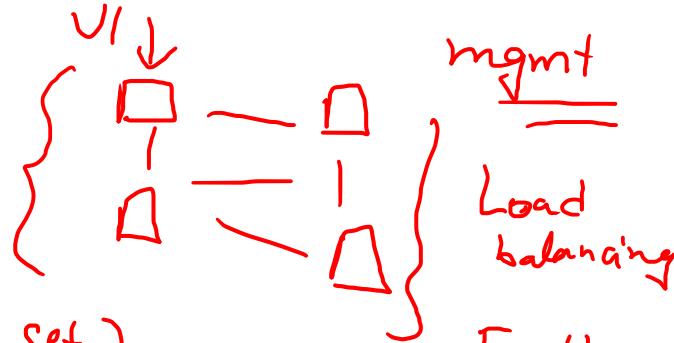
↳ flexible use of n/w
multi-user mgmt →

HPC



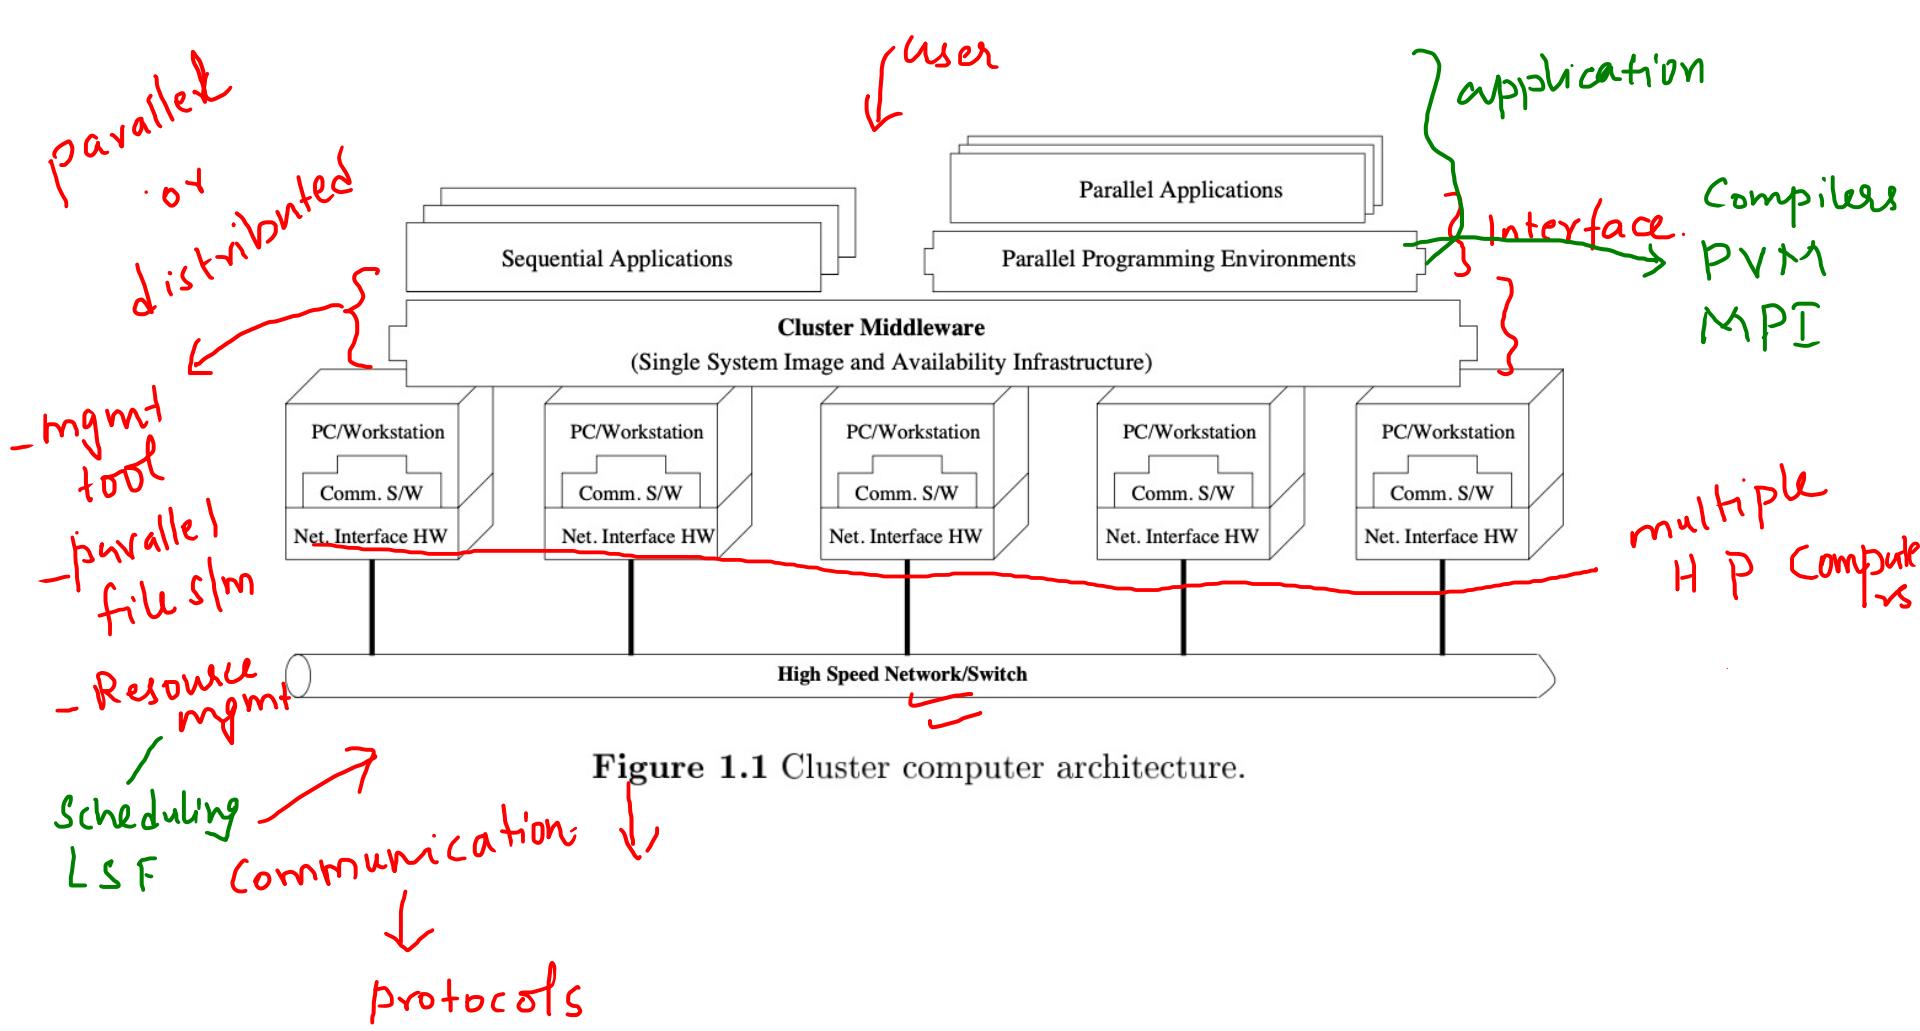
High Performance
computing powers

Allocation

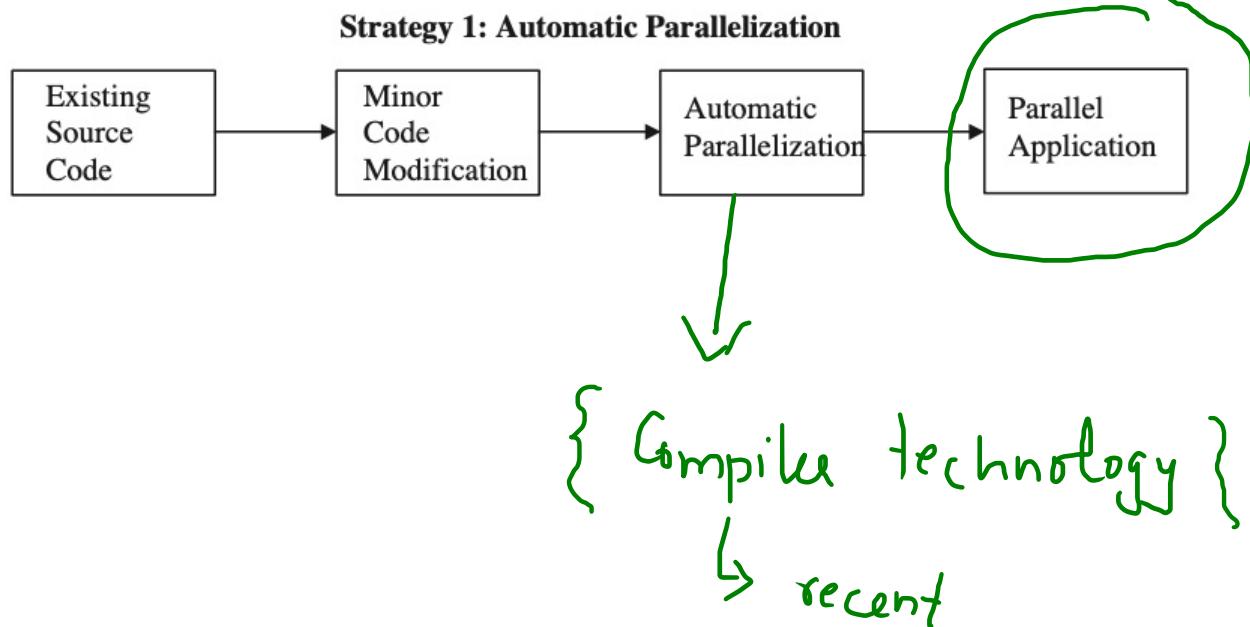


Main motivation

- minimize economic risk
- Consumer off-the-shelf technology

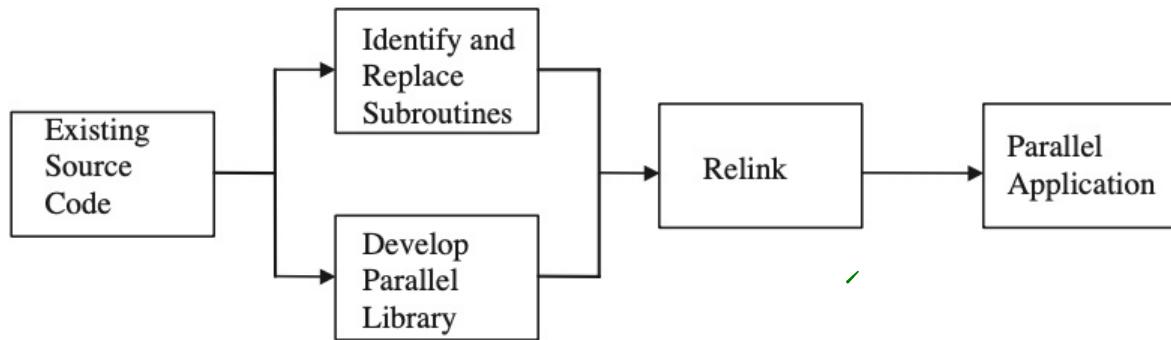


Strategies for developing parallel App^{ms}



Successful

Strategy 2: Parallel Libraries



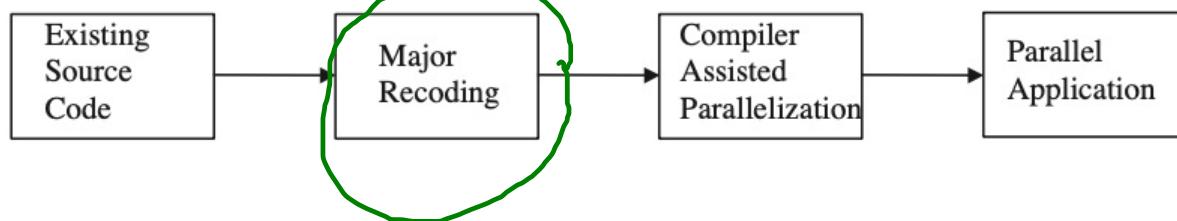
{reusability}

Parallel Libraries

(1) Control structures

(2) provide parallel implⁿ of mathematical routines

Strategy 3: Major Recoding



Scratch

Level of parallelism

h/w

s/w

Code granularity

- 1. Signal } h/w
- 2. Circuit } h/w
- 3. Component } s/w
- 4. S/m level } s/w

(umps of code)
grain size

Code granularity

(latency)

Parallelised by

- (1) Very fine → multiple "inst" issue → Processor
- (2) Fine → data level → Compiler
- (3) Medium → Control level → Programmer
- (4) Large → task - level → "

Level of parallelism

h/w

s/w

Code granularity

- 1. Signal } h/w
- 2. Circuit } h/w
- 3. Component } s/w
- 4. S/m level } s/w

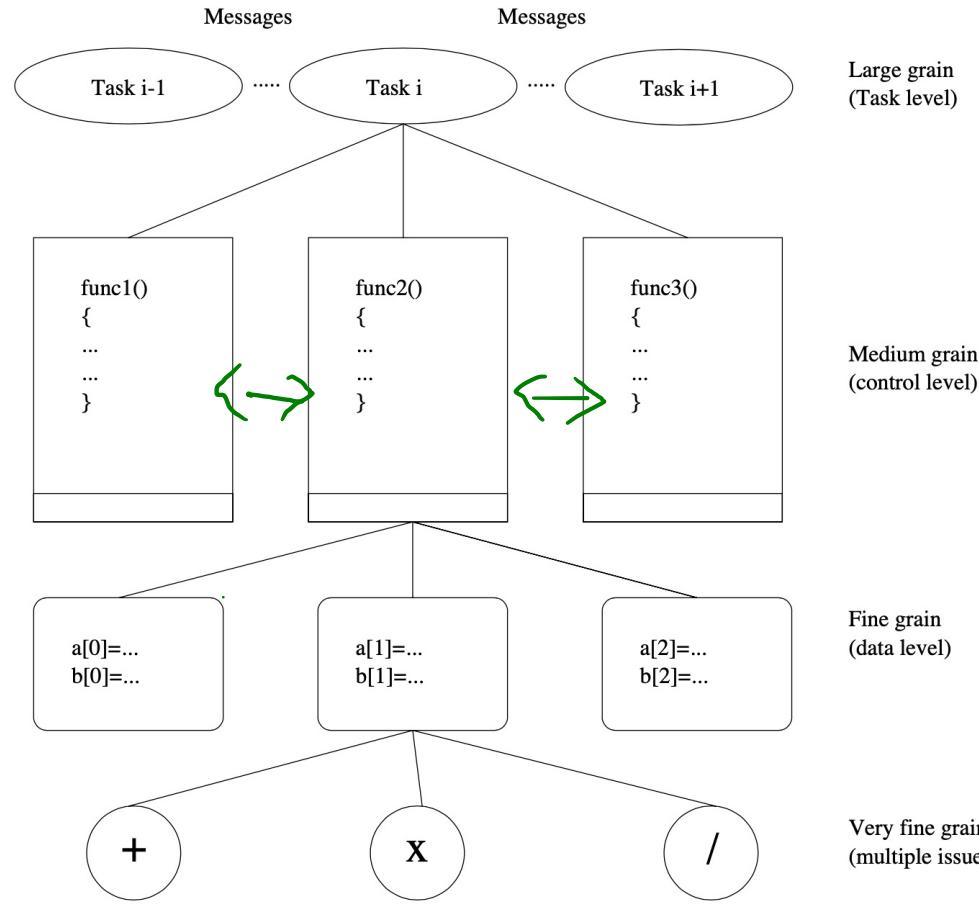
(umps of code)
grain size

Code granularity

(latency)

Parallelised by

- (1) Very fine → multiple "inst" issue → Processor
- (2) Fine → data level → Compiler
- (3) Medium → Control level → Programmer
- (4) Large → task - level → "



Large grain
(Task level)

Medium grain
(control level)

Fine grain
(data level)

Very fine grain
(multiple issue)

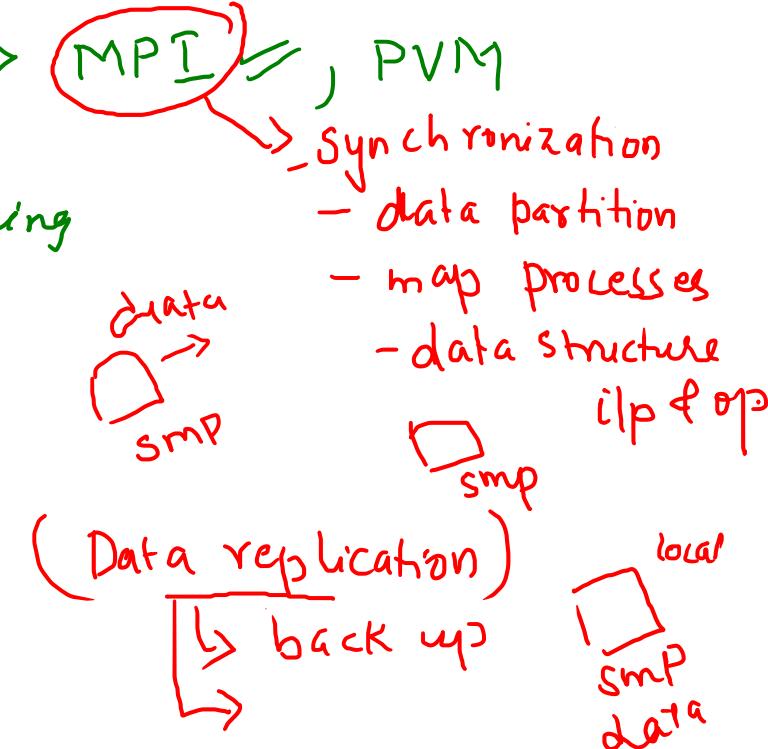
Programmer → Compiler → Processor

Tools

1. Parallelizing Compilers ✓ 3 architectures
2. Parallel languages ✓
3. message passings → MPI, PVM
4. Virtual shared m/y
5. Object Oriented pgmming
6. Pgmming skeleton.

vs implements
shared m/y
pgmming model ←
in a distributed
m/y environment

Distributed Shared m/y
hybrid



Methodical Design of Parallel Algorithms

machine independent ✓ → early

machine-specific tasks. → bit delayed

$$\frac{\# \text{Computation}}{\# \text{Comm}^n}$$

- (1) Partitioning → decomposing
- (2) Communication → Coordinate all the partitioned tasks
- (3) Agglomeration → Evaluate first two stages
- (4) Mapping → Load balancing

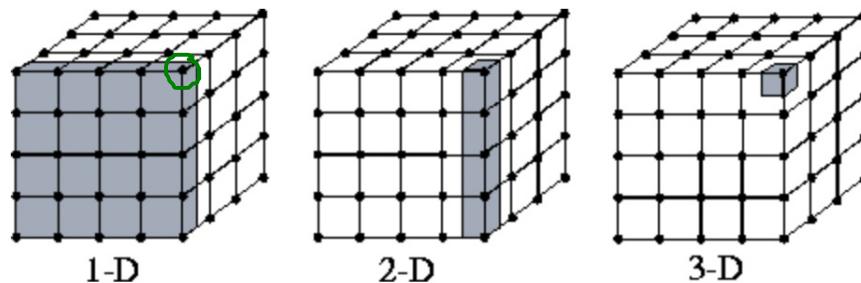
- ↓
(a) maximizing processor utilization
(b) communication cost ↓

{ Scalability $\swarrow h/w$
Efficiency $\swarrow s/w$

Methodical Design of Parallel Algorithms

1. Partitioning

Datastructure

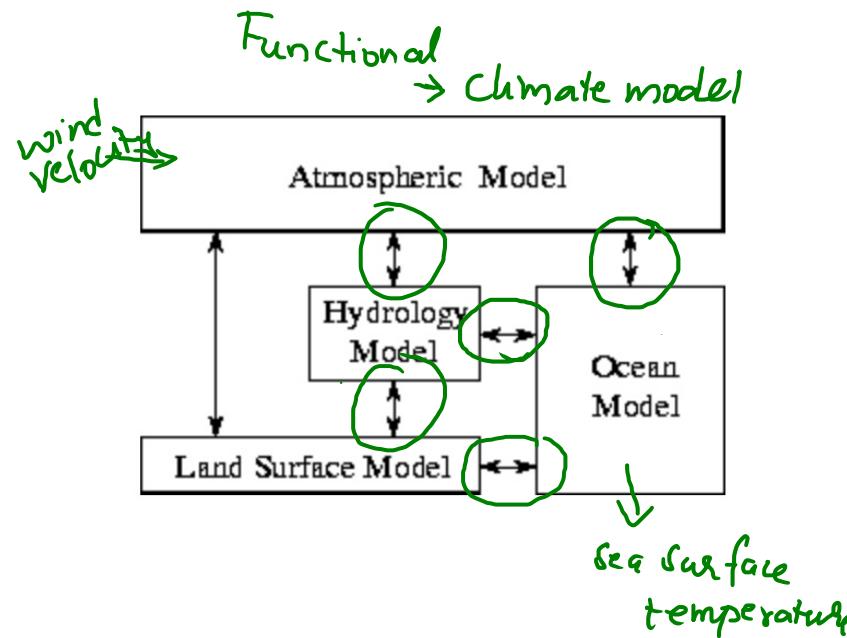


Decomposition → Very fine

Computation data

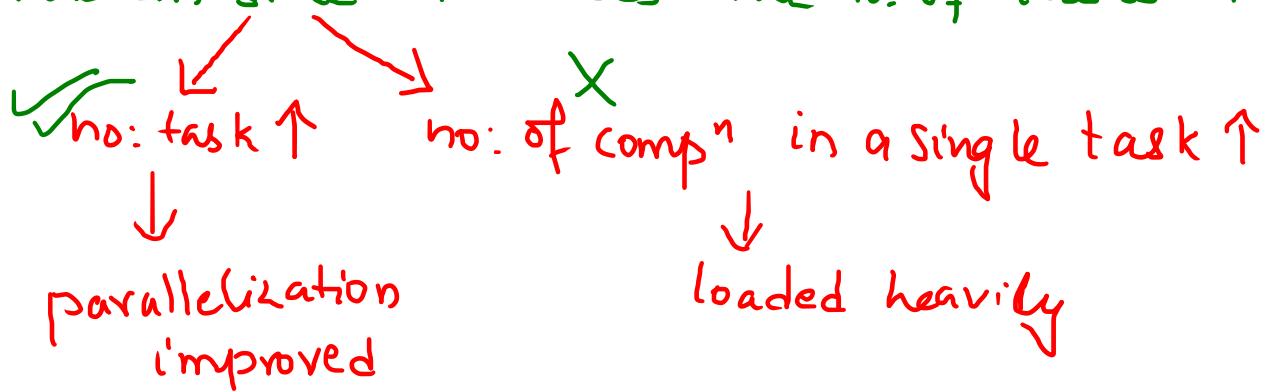
functional decomposition — Computations ← associate data

domain decomposition — first focus is on data — associate computⁿ to data



{ Disjoint sets
Overlapping

- 1. n no. of processors \rightarrow n no. of tasks
- 2. redundant computation / storage reqmts.
- 3. Granularity \rightarrow Equal amount of work to processors
- 4. Problem size \uparrow Does the no. of tasks \uparrow or \downarrow ?

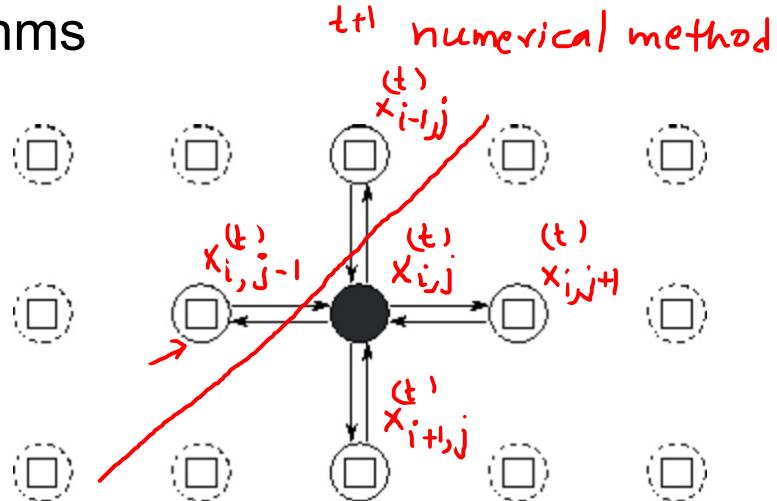


Methodical Design of Parallel Algorithms

Communication

1. channel link
2. messages

$$\frac{\# \text{ computation}}{\# \text{ communication}}$$



grid
 tree
 set of tasks
 prod cons

1. local/global commⁿ ← graph
 2. Structured/unstructured commⁿ
 3. Static/dynamic commⁿ ← varying over time.
 4. Synchronous/asynchronous commⁿ
 ↓
 Consumers

$$x_{i,j}^{(t+1)} = \underbrace{4x_{i,j}^{(t)}}_{\text{Consumed}} + x_{i-1,j}^{(t)} + x_{i+1,j}^{(t)} + x_{i,j+1}^{(t)} + x_{i,j-1}^{(t)}$$

Domain

{ Data disjoint
 Set }
 (Challenging)

Functional

↓
 data flow
 b/w tasks

$\{x_{ij}^{(1)}, x_{ij}^{(2)}, x_{ij}^{(3)} \dots\}$

$x_{i-1,j}^{(1)}, x_{i-1,j}^{(2)} \dots$

$x_{i+1,j}^{(1)} \dots$

$x_{i,j-1}^{(1)} \dots$

$x_{i,j+1}^{(1)} \dots$

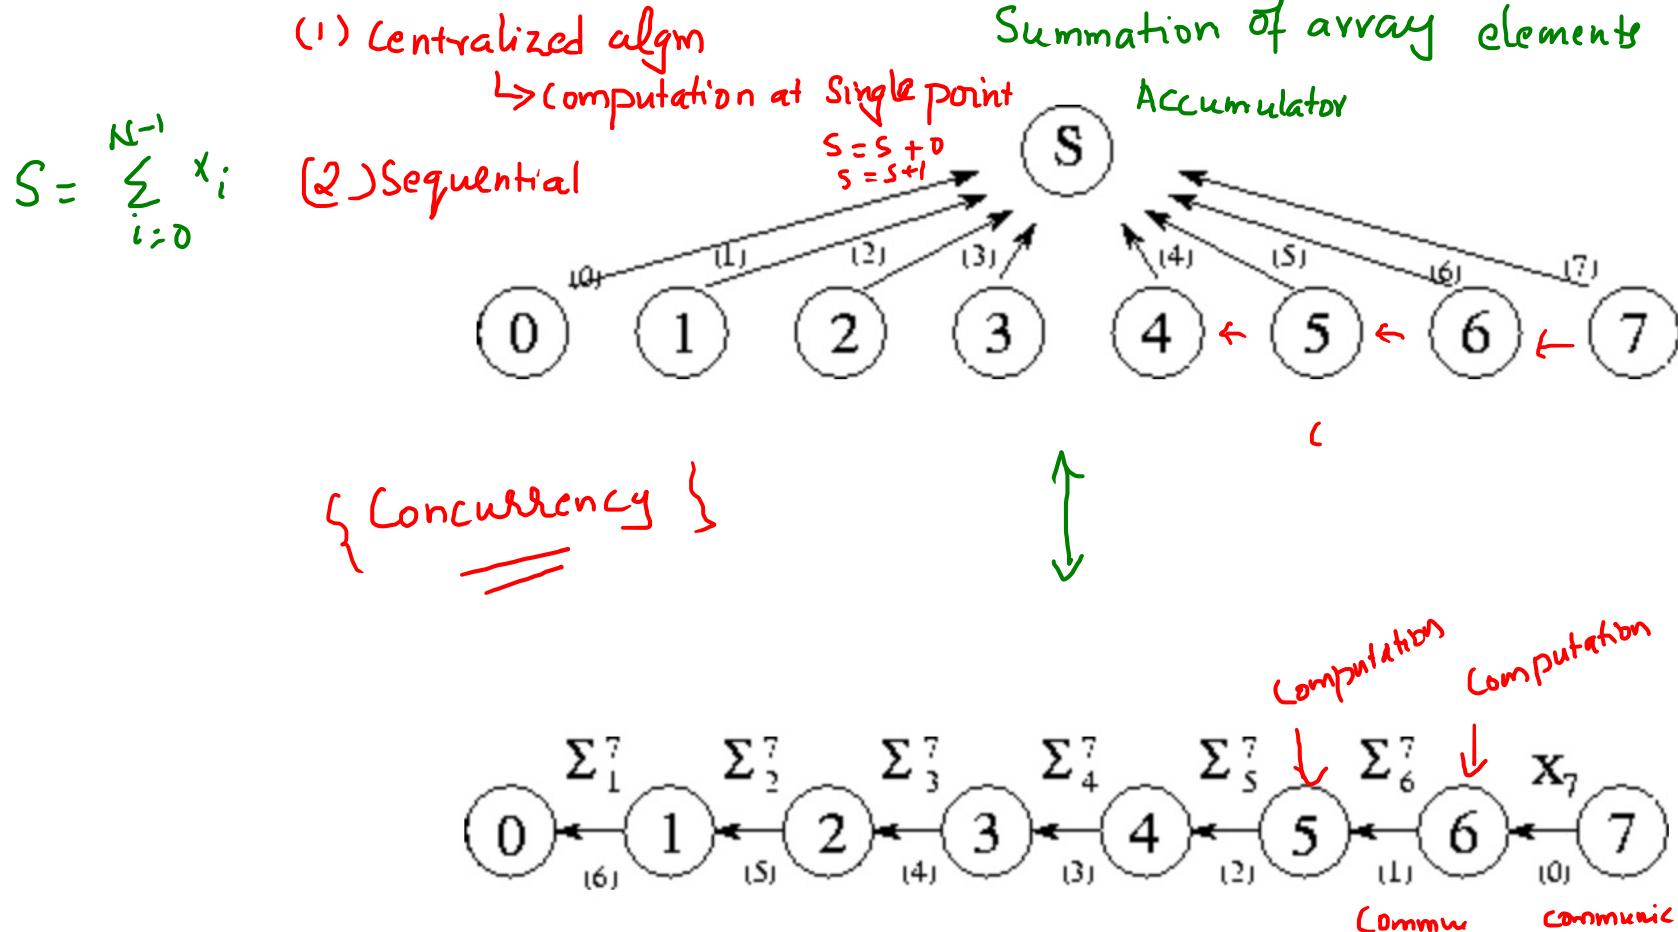
for $t=0$ to $T-1$
 { send $x_{ij}^{(t)}$ to each neighbour
 receive 4 values from
 neighbours
 Compute $x_{ij}^{(t+1)}$

}

x	x	x	x
x	x	x	x
x	x	•	x
-			

x		x
	x	
x	•	x
-		

Methodical Design of Parallel Algorithms



Methodical Design of Parallel Algorithms

Divide & Conquer { $\log n$ }
concurrently

procedure DAC

{

if base case

solve the problem

else

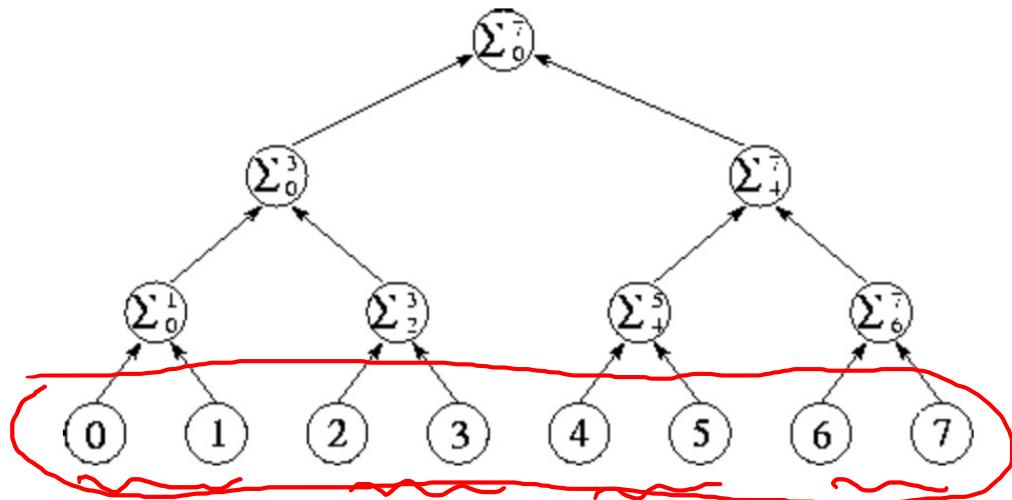
partitioning problem into L & R

Solve L using DAC

Solve R using DAC

Combine L & R to get solution

}

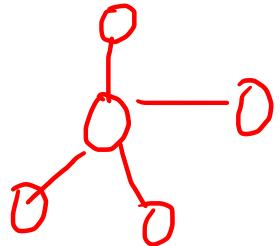


Unstructured

—

graphs

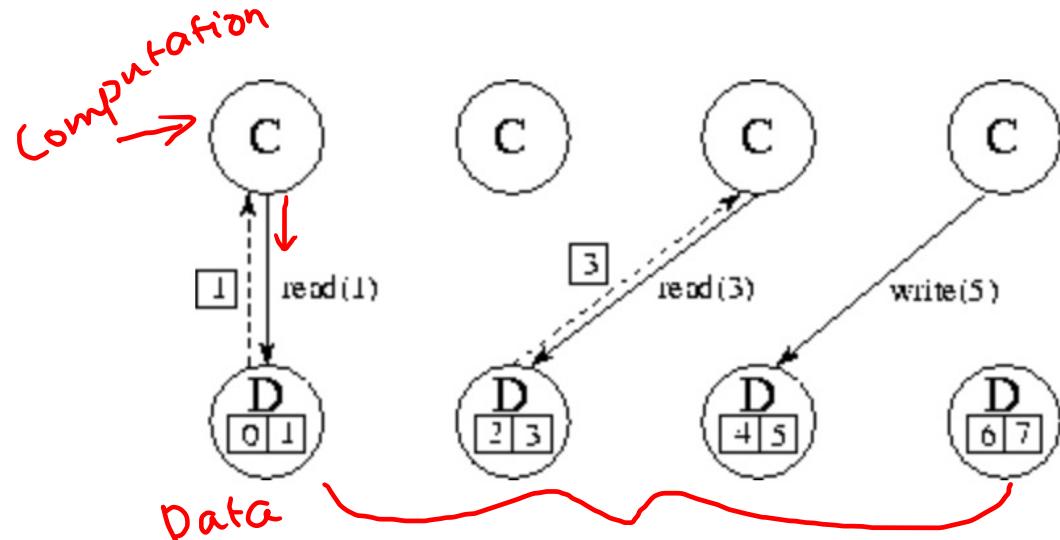
mesh



Dynamic

Methodical Design of Parallel Algorithms

Asynchronous
or
synchronous.



1. polls polling

2.

- (C)
1. $A[i]$
 2. $A[i] \times B[i]$
 3. \rightarrow request C polling
 4. $B[j] \times C[j]$
 5. :
 6. request D polling

Communication —

1. Unbalanced → Scalability ?

2. global →
local

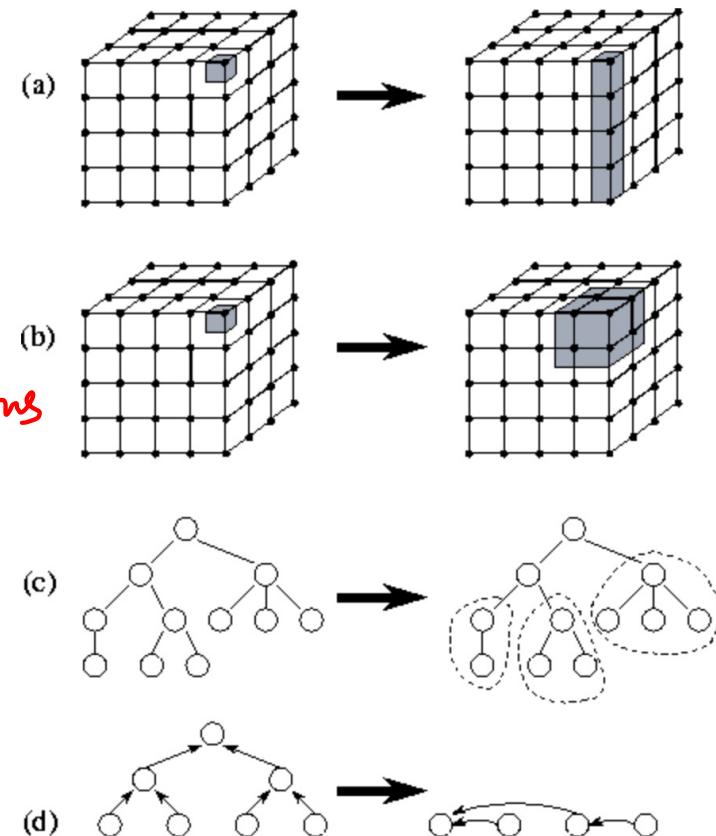
3. efficiency → Concurrency
PAC

Methodical Design of Parallel Algorithms — Agglomeration

Determine → replicate the data

1. reduction in communication cost
increasing computation
2. retain the flexibility w.r.t
Scalability and mapping decisions
3. Reduce the S/w Engineering Cost

task < processors



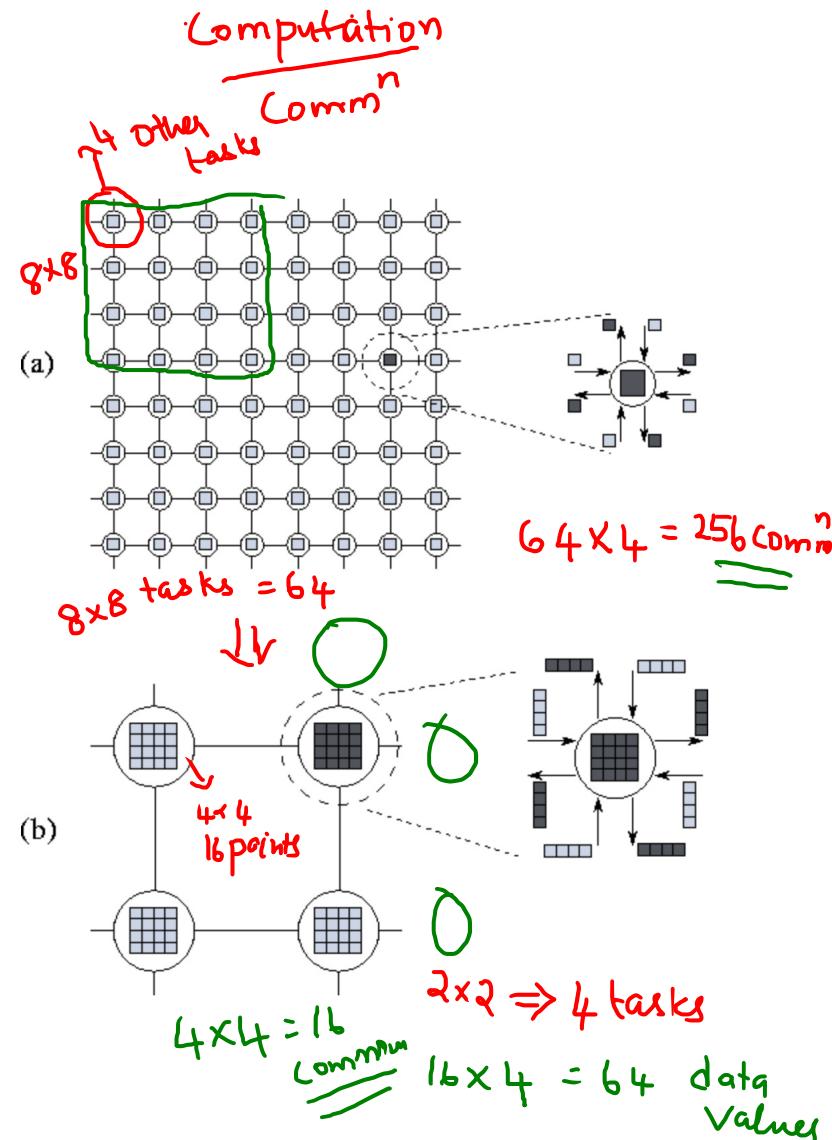
Methodical Design of Parallel Algorithms

1. Increasing Granularity

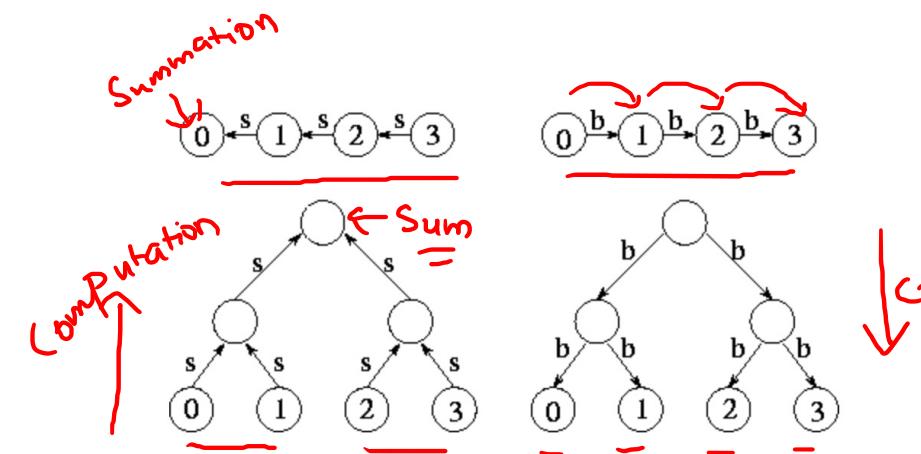
Surface - to - volume effects

Replicating Computations

① ② ③



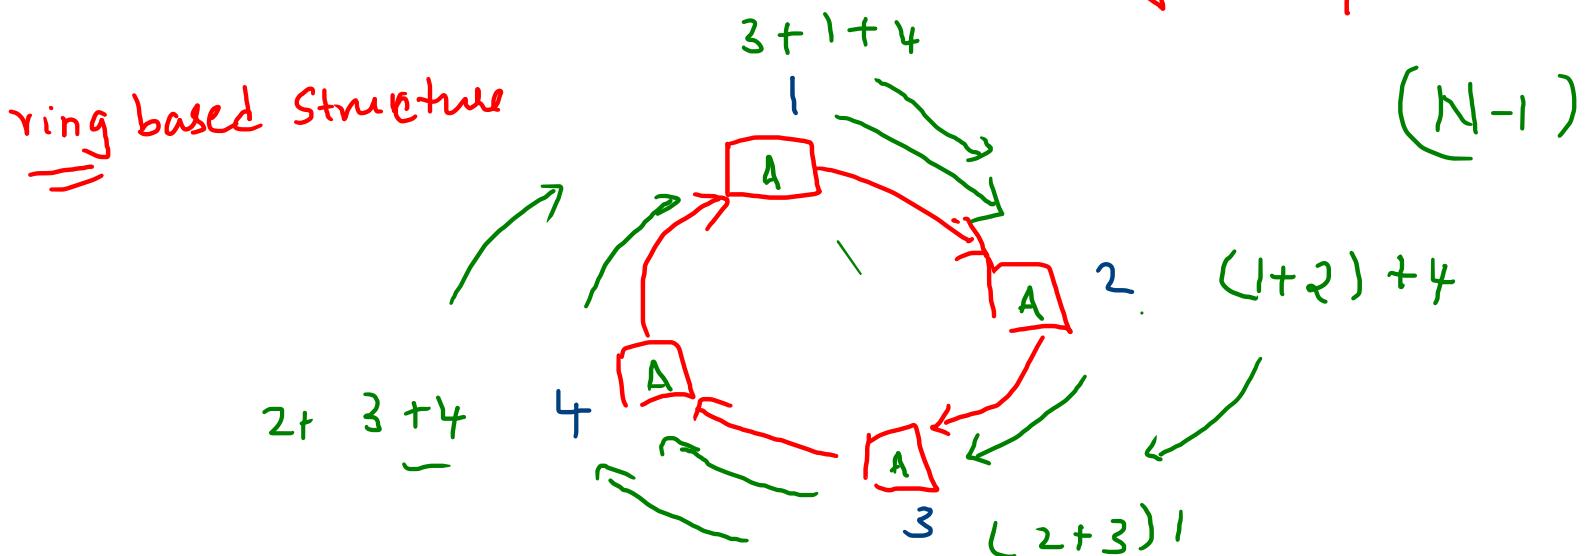
Methodical Design of Parallel Algorithms



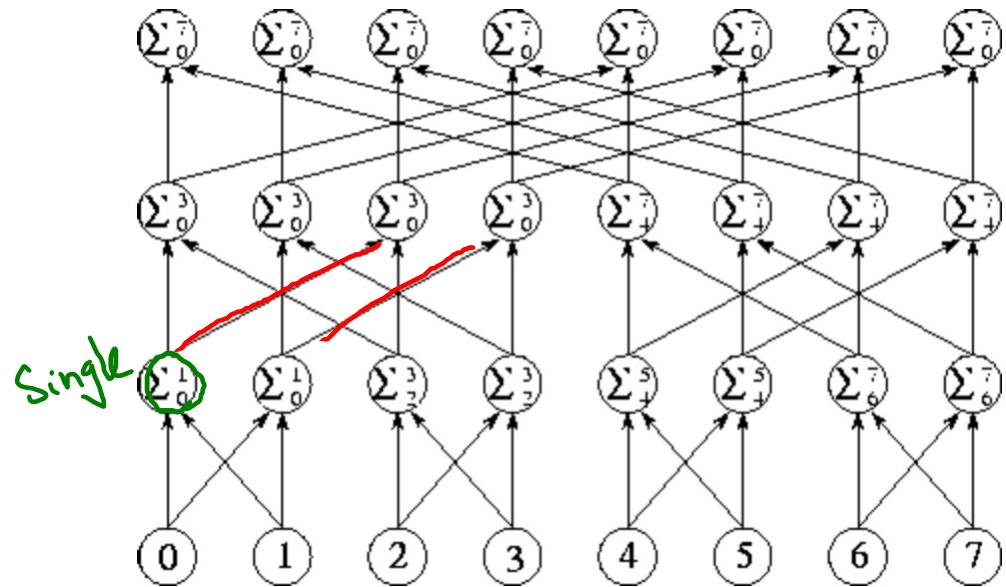
Sum → replicate

S - Summation
 b - broad casting

- $2(N-1)$ steps
- $2 \log N$ steps.



Butterfly Comm



n tasks \Rightarrow single processor

more processors

Waiting
for remote data

{ Overlapping
Compⁿ & Commⁿ }

\geq more no: of tasks than processors

{ mapping strategies }

↳ load balance

4 - Mapping

mapping algms → minimize the total execution time

Goal

- (1) To execute concurrently on different processors.
- (2) Tasks which communicate frequently on the same processor , which increases locality.

Load balancing algms

Dynamic load balancing strategy - Agglomeration & mapping

Probabilistic load balancing

{ local algorithms }

{ Global computation state }

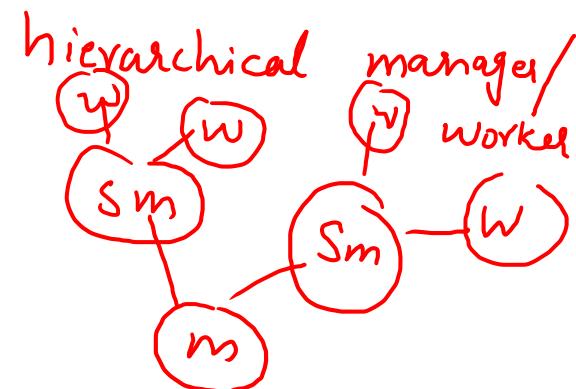
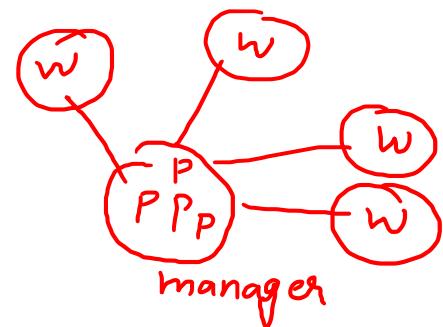
{ domain decomposition }

master / slave

Functional
decomposition

} → task - scheduling Algm →

Decentralized
—
RR



RAM model

Count the primitive operations

*

frequency count

-

$O(n)$

=

$s = \text{sum}$

$\Theta(n)$

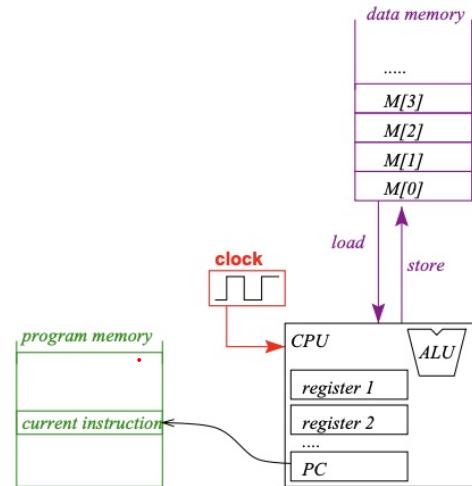
=

$s = d(0)$

do $i = 1, N-1$

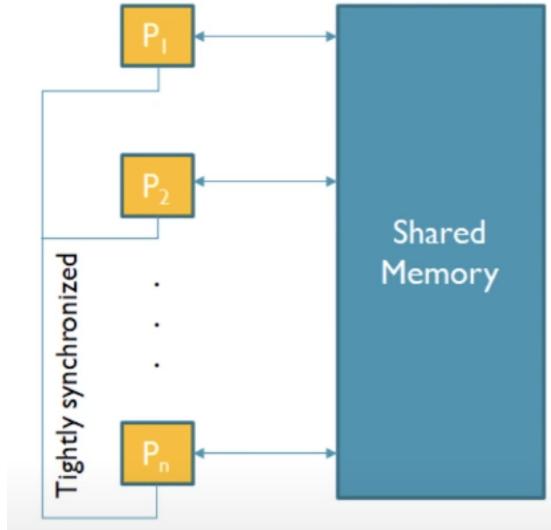
$s = s + d(i)$

end do



PRAM

1. P processors connected to single shared m/y
2. processor id ($1 \leq i \leq p$)
3. SIMD



multiprocessor system with shared memory, on which all n processors operate in lockstep in memory access and program execution operations. Each processor can access any memory location in unit time

PRAM Model Variants

1. EREW → Exclusive Read (Upper bound)

↳ Exclusive Write

↳ does not allow any kind of simultaneous to a single memory location

2. CREW → Concurrent Read exclusive Write

3. ERCW → Exclusive Read Concurrent Write.

4. CRCW — Concurrent Read & Write. (Lower bound)

Writes	Reads	
	Exclusive	Concurrent
Exclusive	EREW	CREW
	Least effective Most simple	Most common
Concurrent	ERCW	CRCW
	Practically no use	Most powerful Most complex

CRCW → How to resolve concurrent writes?

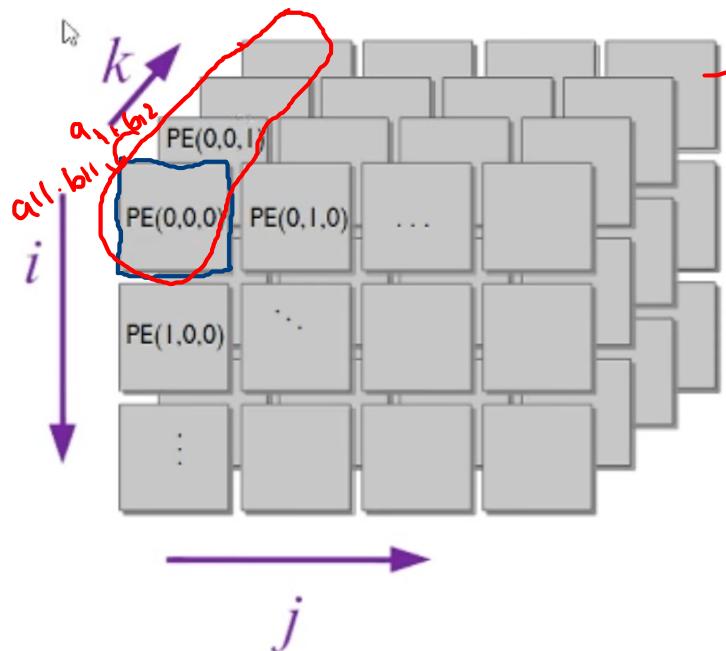
→ (1) Common CRCW → Concurrent writes are allowed when processors are attempting to write the same value

(2) Arbitrary CRCW → a value arbitrarily chosen from the values written to the common my location

(3) Priority CRCW →

(4) Combining CRCW →

Matrix multiplication — PRAM



$$A = \begin{pmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \cdots & a_{mn} \end{pmatrix}, \quad B = \begin{pmatrix} b_{11} & b_{12} & \cdots & b_{1p} \\ b_{21} & b_{22} & \cdots & b_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ b_{n1} & b_{n2} & \cdots & b_{np} \end{pmatrix}$$

$$C = \begin{pmatrix} a_{11}b_{11} + \cdots + a_{1n}b_{n1} & a_{11}b_{12} + \cdots + a_{1n}b_{n2} & \cdots & a_{11}b_{1p} + \cdots + a_{1n}b_{np} \\ a_{21}b_{11} + \cdots + a_{2n}b_{n1} & a_{21}b_{12} + \cdots + a_{2n}b_{n2} & \cdots & a_{21}b_{1p} + \cdots + a_{2n}b_{np} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1}b_{11} + \cdots + a_{mn}b_{n1} & a_{m1}b_{12} + \cdots + a_{mn}b_{n2} & \cdots & a_{m1}b_{1p} + \cdots + a_{mn}b_{np} \end{pmatrix}$$

→ individual multiplication

SIMD

$$\{ C_{ij} = A_{ik} K \times B_{ki} \}$$

$$\downarrow O(1)$$

$$\downarrow O(n^3)$$

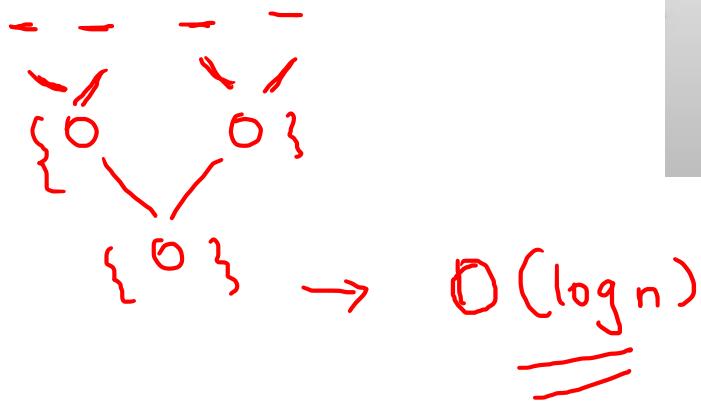
↓
Strassen's mult

$$\downarrow O(n^{2.81})$$

Individual processors

Step 1

1. Read $A(i, k)$
2. Read $B(k, j)$
3. Compute $A(i, k) \times B(k, j)$
4. Store in $C(i, j, k)$



Step 2

```
1.  $\ell \leftarrow n$ 
2. Repeat
    $\ell \leftarrow \ell/2$ 
   if ( $k < \ell$ ) then
      begin
         Read  $C(i, j, k)$ 
         Read  $C(i, j, k + \ell)$ 
         Compute  $C(i, j, k) + C(i, j, k + \ell)$ 
         Store in  $C(i, j, k)$ 
      end
   until ( $\ell = 1$ )
```

Summation

of

n numbers

\downarrow
DAC

$\underline{\underline{O(n^3)}}$

Merging 2 sorted list

Parallel Merge

L_1 L_2
 $n/2$ $n/2$ $\rightarrow n$ elements

$h \rightarrow$ processors

binary Search

$A = \{a_1, a_2, \dots\}$
 $B = \{b_1, b_2, \dots\}$
 $C = \{a_1, a_2, b_1, \dots\}$

$O(n)$

$A[1]$ P_3
 $A = 1, 5, \boxed{7}, 13, 17, 19, 23 \rightarrow \text{lower}$

- Processors
One for each element

$B = 2, \textcircled{4}, 8, 11, 12, 21, 24 \rightarrow \text{Upper}$
 $A[9]$ $A[16]$

$C = ?$

Task of P_3

$(i-1) + (\text{high} - n/2) \Rightarrow$ position of 7 in merged list.

Binary Search : $A[i=3]$

↓

$A[3]$

$i-1 = (3-1) = 2$ elements in the lower array

7 is larger 2 elements "

high = index of the largest integer smaller than 7

index = 10 $\uparrow (\text{high} - n/2) = 10 - 8 = 2$ elements

in upper array

$A = [1, 5, 7, 13, 17, 19, 23]$ → lower
 $B = [2, 4, 8, 11, 12, 21, 24]$ → upper
 $C = ?$

- Processors
 One for each element

Task of P_{11}

$$(i - (n/2 + 1) + \text{high})$$

$$A[i=11] = 8$$

$$i - (n/2 + 1) = 11 - (8 + 1) = 2$$

Binary Search with $A[11]$ in the lower array
 $\text{high} = \text{index of largest integer smaller than } 8$
 $\text{high} = 3$

$A = [1, 5, 7, 13, 17, 19, 23]$ → lower
 $B = [2, 4, 8, 11, 12, 21, 24]$ → upper
 $C = ?$

- Processors
 One for each element

Task of P_{11}

$$(i - (n/2 + 1) + \text{high})$$

$$A[i=11] = 8$$

$$i - (n/2 + 1) = 11 - (8 + 1) = 2$$

Binary Search with $A[11]$ in the lower array
 $\text{high} = \text{index of largest integer smaller than } 8$
 $\text{high} = 3$

MERGE.LISTS (CREW PRAM):

Given: Two sorted lists of $n/2$ elements each, stored in $A[1] \dots A[n/2]$ and $A[(n/2) + 1] \dots A[n]$
 The two lists and their unions have disjoint values
 Final condition: Merged list in locations $A[1] \dots A[n]$
 Global $A[1 \dots n]$
 Local $x, low, high, index$
 begin
 spawn (P_1, P_2, \dots, P_n)
 for all P_i where $1 \leq i \leq n$ do
 (Each processor sets bounds for binary search)
 if $i \leq n/2$ then
 $low \leftarrow (n/2) + 1$
 $high \leftarrow n$
 else
 $low \leftarrow 1$
 $high \leftarrow n/2$
 endif

↑↑
Array C

Binary search

(Each processor performs binary search) → ? T/F

```

x ← A[i]
repeat
    index ← ⌊(low + high)/2⌋
    if x < A[index] then
        high ← index - 1
    else
        low ← index + 1
    endif
    until low > high
    ( Put value in correct position on merged list )
    A[high + i - n/2] ← x
endfor
end

```

{ lower array }
 { upper array }

Sequential algm

$O(n)$

Parallel Merge
 Time Complexity ?

Binary S = $O(\log n)$

Overall TC = $O(n \log n)$

$(\text{const} + B)$ → one element

MERGE.LISTS (CREW PRAM):

Given: Two sorted lists of $n/2$ elements each, stored in $A[1] \dots A[n/2]$ and $A[(n/2) + 1] \dots A[n]$
 The two lists and their unions have disjoint values
 Final condition: Merged list in locations $A[1] \dots A[n]$

```

Global  A[1 ... n]
Local   x, low, high, index
begin
  spawn (P1, P2, ..., Pn)
  for all Pi where 1 ≤ i ≤ n do
    { Each processor sets bounds for binary search }
    if i ≤ n/2 then
      low ← (n/2) + 1
      high ← n
    else
      low ← 1
      high ← n/2
    endif
  endfor
end
    
```

↑↑
Array C

Binary search

{ Each processor performs binary search } → ? T/F

```

x ← A[i]
repeat
  index ← ⌊(low + high)/2⌋
  if x < A[index] then
    high ← index - 1
  else
    low ← index + 1
  endif
  until low > high
  { Put value in correct position on merged list }
  A[high + i - n/2] ← x
endfor
end
    
```

{ lower array }
 { upper array }

Sequential algm

O(n)

Parallel Merge

Time Complexity ?

(Const + B) → One element

Binary S = O(log n)

Overall TC = O(n log n)

Work - Depth Model

→ DAG

Single processor

Time taken \propto number of operations

$T_1 \rightarrow$ amount of time algm take on
one processor

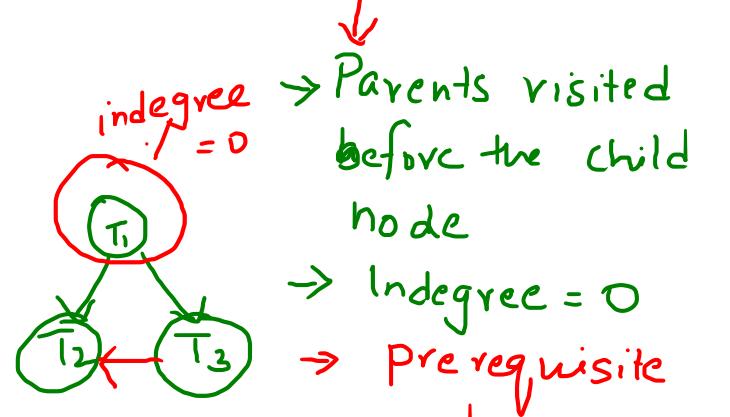
$T_p \rightarrow$ " P processors.

Wait for the slowest processor to finish

This is known as the depth of an algm

Work :- no: of computations \times no: of processors used

↳ Topological ordering



course

Lower bound

$$\frac{T_1}{p} \leq T_p$$

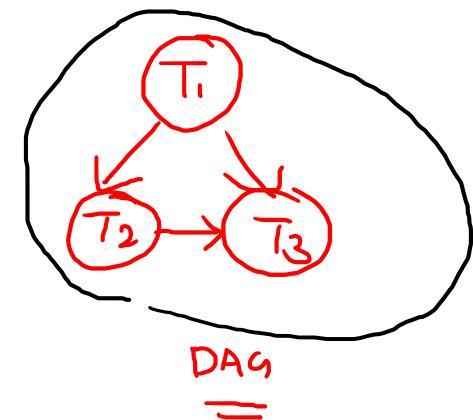
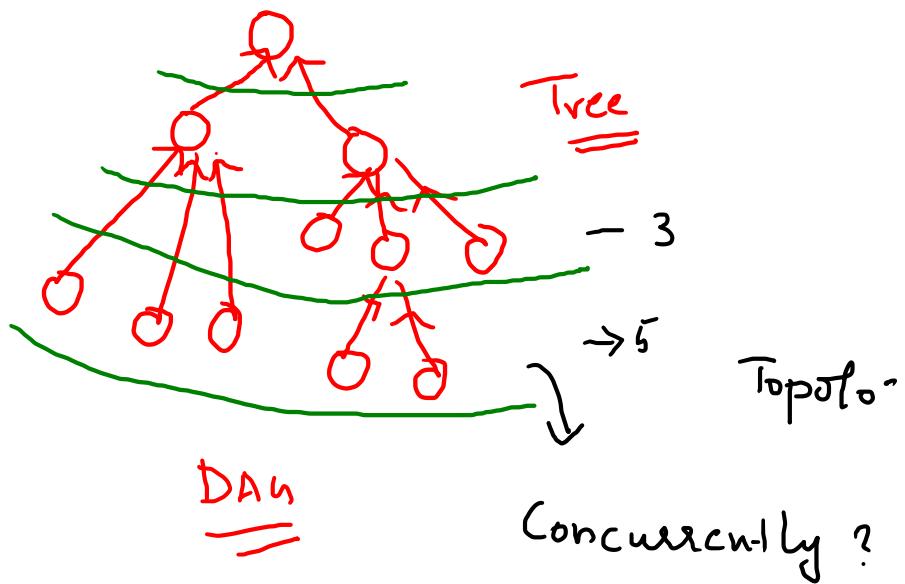
Depth of PRAM with Infinite Processors

- Representing algms as DAG's

↳ multiple computations vs infinite processors

↳ representing the dependences b/w operations
in an algm using a DAG.

Constructing a DAG from an algm.



depth = T_{∞} = depth of computation DAG.